

Dongfu Jiang

PhD Student in Computer Science, University of Waterloo | Waterloo, Canada | Updated March 20, 2026

✉ dongfu.jiang@uwaterloo.ca | 🏠 jdf-prog.github.io | 🔗 [jdf-prog](#) | 🐦 [@DongfuJiang](#)

Education

University of Waterloo

PhD in Computer Science

Advisor: Prof. Wenhua Chen

Waterloo, Canada

Sep 2023 – Sep 2026 (expected)

Zhejiang University

BE in Computer Science

Advisor: Prof. Zhou Zhao | GPA: 3.97/4.00

Hangzhou, China

Sep 2019 – Jun 2023

Honors and Awards

2025

🏆 Outstanding Paper Award, TMLR 2025, for *MANTIS* (1 / 1539 selected) 2025

🏆 Journal to Conference Certificate, TMLR 2025, for *StructEval* 2025

2024

🏆 Best Paper Finalist and Oral, CVPR 2024, for *MMMU* (24 / 11,532 selected) 2024

2022

Xiaomi Scholarship (Zhejiang University) 2022

2021

Zhejiang University First-Class Scholarship (top 3%) 2021

2020

National Scholarship, Zhejiang University (top 1%) 2020

Zhejiang University First-Class Scholarship (top 3%) 2020

Research Experience

Research Intern, NVIDIA ADLR

Advisors: Ping Wei, Mohammad Shoeybi

Agentic reinforcement learning for tool use; Contributing to post-training of nemotron family of models

Santa Clara, US

Aug 2025 – Present

Research Intern, Allen Institute for AI

Advisors: (Bill) Yuchen Lin, Qianhui Qin

Active learning with verbalized human feedback

Seattle, US

Jun 2024 – Sep 2024

Research Associate, SeaAI

Advisors: Qian Liu, Tianyu Pang, Chao Du

MANTIS: Interleaved multi-image instruction tuning

Singapore (remote)

Feb 2024 – Sep 2025

Research Intern, University of Southern California

Advisors: Xiang Ren, (Bill) Yuchen Lin

LLM-Blender: Ensembling large language models with pairwise ranking and generative fusion

US (remote)

Mar 2022 – Mar 2023

Publications

* denotes equal contribution where marked. [Google Scholar](#); [Semantic Scholar](#)

Manuscripts and Preprints

2026

[1] Nemotron-Cascade 2: Post-Training LLMs with Cascade RL and Multi-Domain On-Policy Distillation

Zhuolin Yang, Zihan Liu, Yang Chen, Wenliang Dai, Boxin Wang, Sheng-Chieh Lin, Chankyu Lee, Yangyi Chen, **Dongfu Jiang**, Jiafan He, Renjie Pi, Grace Lam, Nayeon Lee, Alexander Bukharin, Mohammad Shoeybi, Bryan Catanzaro, Wei Ping

arXiv preprint | [arXiv](#)

- [2] **Nemotron 3 Super: Open, Efficient Mixture-of-Experts Hybrid Mamba-Transformer Model for Agentic Reasoning**
Core Contributor
Technical report, March 11, 2026 | [Report](#)
- [3] **OpenResearcher: A Fully Open Pipeline for Long-Horizon Deep Research Trajectory Synthesis**
Zhuofeng Li, **Dongfu Jiang***, Xueguang Ma, Haoxiang Zhang, Ping Nie, Yuyu Zhang, Kai Zou, Jianwen Xie, Yu Zhang, Wenhui Chen
Blog post, Feb 9, 2026 | [Blog post](#) | [TIGER-AI-Lab/OpenResearcher](#) ★ 420
- 2025
- [4] **VideoScore2: Think before You Score in Generative Video Evaluation**
Xuan He, **Dongfu Jiang***, Ping Nie, Minghao Liu, Zhengxuan Jiang, Mingyi Su, Wentao Ma, Junru Lin, Chun Ye, Yi Lu, Keming Wu, Benjamin Schneider, Quy Duc Do, Zhuofeng Li, Yiming Jia, Yuxuan Zhang, Guo Cheng, Haozhe Wang, Wangchunshu Zhou, Qunshu Lin, Yuanxing Zhang, Ge Zhang, Wenhao Huang, Wenhui Chen
arXiv preprint | [TIGER-AI-Lab/VideoScore2](#) ★ 18
- [5] **VerlTool: Towards Holistic Agentic Reinforcement Learning with Tool Use**
Dongfu Jiang, Yi Lu, Zhuofeng Li, Zhiheng Lyu, Ping Nie, Haozhe Wang, Alex Su, Hui Chen, Kai Zou, Chao Du, Tianyu Pang, Wenhui Chen
arXiv preprint | [TIGER-AI-Lab/verl-tool](#) ★ 872
- Peer-reviewed Publications**
- 2026
- [6] **Critique-Coder: Enhancing Coder Models by Critique Reinforcement Learning**
Chi Ruan, **Dongfu Jiang**, Yubo Wang, Wenhui Chen
ICLR 2026 | [TIGER-AI-Lab/Critique-Coder](#) ★ 11
- 2025
- [7] **StructEval: Benchmarking LLMs' Capabilities to Generate Structural Outputs**
Jialin Yang*, **Dongfu Jiang***, Lipeng He, Sherman Siu, Yuxuan Zhang, Disen Liao, Zhuofeng Li, Huaye Zeng, Yiming Jia, Haozhe Wang, Benjamin Schneider, Chi Ruan, Wentao Ma, Zhiheng Lyu, Yifei Wang, Yi Lu, Quy Duc Do, Ziyang Jiang, Ping Nie, Wenhui Chen
TMLR 2025 (Journal to Conference Certificate) | [TIGER-AI-Lab/StructEval](#) ★ 13
- [8] **QuickVideo: Real-Time Long Video Understanding with System Algorithm Co-Design**
Benjamin Schneider*, **Dongfu Jiang***, Chao Du, Tianyu Pang, Wenhui Chen
TMLR 2025 | [TIGER-AI-Lab/QuickVideo](#) ★ 55
- [9] **General-Reasoner: Advancing LLM Reasoning Across All Domains**
Xueguang Ma*, Qian Liu*, **Dongfu Jiang**, Ge Zhang, Zejun Ma, Wenhui Chen
NeurIPS 2025 | [TIGER-AI-Lab/General-Reasoner](#) ★ 195
- [10] **ACECODER: Acing Coder RL via Automated Test-Case Synthesis**
Huaye Zeng*, **Dongfu Jiang***, Haozhe Wang, Ping Nie, Xiaotong Chen, Wenhui Chen
ACL 2025 | [TIGER-AI-Lab/AceCoder](#) ★ 93
- [11] **MEGA-Bench: Scaling Multimodal Evaluation to over 500 Real-World Tasks**
Jiacheng Chen*, Tianhao Liang*, Sherman Siu*, Zhengqing Wang, Kai Wang, Yubo Wang, Yuansheng Ni, Wang Zhu, Ziyang Jiang, Bohan Lyu, **Dongfu Jiang**, Xuan He, Yuan Liu, Hexiang Hu, Xiang Yue, Wenhui Chen
ICLR 2025 | [TIGER-AI-Lab/MEGA-Bench](#) ★ 77
- 2024
- [12] **GenAI Arena: An Open Evaluation Platform for Generative Models**
Dongfu Jiang*, Max W.F. Ku*, Tianle Li*, Yuansheng Ni, Shizhuo Sun, Rongqi Fan, Wenhui Chen
NeurIPS 2024 Dataset and Benchmark Track
- [13] **VideoScore: Building Automatic Metrics to Simulate Fine-grained Human Feedback for Video Generation**
Xuan He*, **Dongfu Jiang***, Ge Zhang, Max Ku, Achint Soni, Sherman Siu, Haonan Chen, Abhranil Chandra, Ziyang Jiang, Aaran Arulraj, Kai Wang, Quy Duc Do, Yuansheng Ni, Bohan Lyu, Yaswanth Narsupalli, Rongqi Fan, Zhiheng Lyu, Yuchen Lin, Wenhui Chen
EMNLP 2024 | [TIGER-AI-Lab/VideoScore](#) ★ 103
- [14] **WildVision: Evaluating Vision-Language Models in the Wild with Human Preferences**
Yujie Lu, **Dongfu Jiang**, Wenhui Chen, William Yang Wang, Yejin Choi, Bill Yuchen Lin
NeurIPS 2024 Dataset and Benchmark Track
- [15] **MANTIS: Interleaved Multi-Image Instruction Tuning**
Dongfu Jiang, Xuan He, Huaye Zeng, Cong Wei, Max Ku, Qian Liu, Wenhui Chen
TMLR 2024 | [TIGER-AI-Lab/Mantis](#) ★ 231
- 🏆 **Featured Certificate / Outstanding Paper Award at TMLR 2025 (1 / 1539 selected)**

[16] **VIEScore: Towards Explainable Metrics for Conditional Image Synthesis Evaluation**

Max Ku, **Dongfu Jiang**, Cong Wei, Xiang Yue, Wenhui Chen
ACL 2024 | [TIGER-AI-Lab/VieScore](#) ★ 56

[17] **MMMU: A Massive Multi-discipline Multimodal Understanding and Reasoning Benchmark for Expert AGI**

Xiang Yue, Yuansheng Ni, Kai Zhang, Tianyu Zheng, Ruqi Liu, Ge Zhang, Samuel Stevens, **Dongfu Jiang**, Weiming Ren, Yuxuan Sun, Cong Wei, Botao Yu, Ruibin Yuan, Renliang Sun, Ming Yin, Boyuan Zheng, Zhenzhu Yang, Yibo Liu, Wenhao Huang, Huan Sun, Yu Su, Wenhui Chen
CVPR 2024 | [MMMU-Benchmark/MMMU](#) ★ 512
🏆 **Best Paper Finalist and Oral (24 / 11,532 selected)**

[18] **TIGERScore: Towards Building Explainable Metric for All Text Generation Tasks**

Dongfu Jiang*, Yishan Li*, Ge Zhang, Wenhao Huang, Bill Yuchen Lin, Wenhui Chen
TMLR 2024 | [TIGER-AI-Lab/TIGERScore](#) ★ 31

2023

[19] **LLM-Blender: Ensembling Large Language Models with Pairwise Ranking and Generative Fusion**

Dongfu Jiang, Xiang Ren, Bill Yuchen Lin
ACL 2023 | [yuchenlin/LLM-Blender](#) ★ 969

Selected Impact

- 🔊 First / co-first works received broad online coverage, including *MarkTechPost* features on *LLM-Blender*, *GenAI-Arena*, *OpenResearcher*, and *AceCoder*. 2023–2026
- 🔊 *MMMU* has been cited by major multimodal model and benchmark works including *Llama 3*, *LLaVA-OneVision*, *Cambrian-1*, *Kimi-VL*, and *Video-MME*. 2024–2026
- 🔊 *LLM-Blender* has been cited by representative LLM systems and evaluation works including *FrugalGPT*, *Prometheus 2*, *RewardBench*, *SimPO*, and *Mixture-of-Agents*. 2023–2026
- 🔊 *MANTIS* has been cited by multimodal follow-up works including *LLaVA-OneVision*, *LLaVA-NeXT-Interleave*, *Molmo / PixMo*, *InternVL3*, and *MMMU-Pro*. 2024–2026
- 🔊 *VerlTool* has been cited by later agentic RL and tool-use works including *DeepAgent*, *SkyRL*, *AgentFlow*. 2025–2026
- 🔊 *OpenResearcher*'s data has been adopted by NVIDIA's *Nemotron* family of models. 2026
- 🔊 *AceCoder* synthesized prompts were used in the coding RL data mixture of *OLMo 3*. 2025
- 🔊 *WildVision* has been cited by follow-up multimodal evaluation and alignment works including *LLaVA-Critic*, *InternVL3*, *InternVL3.5*, and *Mammoth-VL*. 2024–2026

Talks

- [2025-09-09] Invited talk at TikTok Singapore on *VerlTool*. [slides](#)
- [2024-02-19] Invited talk at the Vector NLP Workshop about *TIGERScore*. [slides](#)
- [2023-07-27] Invited talk at the AutoGPT Reading Group about *LLM-Blender*. [slides](#)

Paper Reviewing

ACL ARR (2024, 2025, 2026) , ICML (2026) , NeurIPS (2024, 2025, 2026) , ICLR (2024, 2026) , TMLR (2024, 2025, 2026) , AISTATS (2025) , AAAI (2025) , ECCV (2026) .